MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Predicting data over time

Chris Holdgraf
Fellow, Berkeley Institute for Data Science

Classification vs. Regression

CLASSIFICATION

```
classification_model.predict(X_test)
array([0, 1, 1, 0])
```

REGRESSION

```
regression_model.predict(X_test)
array([0.2, 1.4, 3.6, 0.6])
```

Correlation and regression

- Regression is similar to calculating correlation, with some key differences

  - **Regression**: A process that results in a formal model of the data

  - **Correlation**: A statistic that describes the data. Less information than regression model.

Correlation between variables often changes over time

- Timeseries often have patterns that change over time

- Two timeseries that seem correlated at one moment may not remain so over time

Visualizing relationships between timeseries

```python
fig, axs = plt.subplots(1, 2)

# Make a line plot for each timeseries
axs[0].plot(x, c='k', lw=3, alpha=.2)
axs[0].plot(y)
axs[0].set(xlabel='time', title='X values = time')

# Encode time as color in a scatterplot
axs[1].scatter(x_long, y_long, c=np.arange(len(x_long)), cmap='viridis')
axs[1].set(xlabel='x', ylabel='y', title='Color = time')
```

Visualizing two timeseries

Regression models with scikit-learn

```python
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X, y)
model.predict(X)
```

Visualize predictions with scikit-learn

```python
alphas = [.1, 1e2, 1e3]
ax.plot(y_test, color='k', alpha=.3, lw=3)
for ii, alpha in enumerate(alphas):
    y_predicted = Ridge(alpha=alpha).fit(X_train, y_train).predict(X_test)
    ax.plot(y_predicted, c=cmap(ii / len(alphas)))
ax.legend(['True values', 'Model 1', 'Model 2', 'Model 3'])
ax.set(xlabel="Time")
```

Visualize predictions with scikit-learn

Scoring regression models

- Two most common methods:

    - Correlation (r)

    - Coefficient of Determination (R^2)

Coefficient of Determination (R^2)

- The value of R^2 is bounded on the top by 1, and can be infinitely low

- Values closer to 1 mean the model does a better job of predicting outputs

$$1 - \frac{error(model)}{variance(testdata)}$$

R^2 in scikit-learn

```python
from sklearn.metrics import r2_score
print(r2_score(y_predicted, y_test))
0.08
```

MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Let's practice!

MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Cleaning and improving your data

Chris Holdgraf
Fellow, Berkeley Institute for Data Science

Data is messy

- Real-world data is often messy

- The two most common problems are *missing data* and *outliers*

- This often happens because of human error, machine sensor malfunction, database failures, etc

- Visualizing your raw data makes it easier to spot these problems

What messy data looks like

Interpolation: using time to fill in missing data

- A common way to deal with missing data is to *interpolate* missing values

- With timeseries data, you can use time to assist in interpolation.

- In this case, **interpolation** means using using the *known* values on either side of

  a gap in the data to make assumptions about what's missing.

Interpolation in Pandas

```python
# Return a boolean that notes where missing values are
missing = prices.isna()

# Interpolate linearly within missing windows
prices_interp = prices.interpolate('linear')

# Plot the interpolated data in red and the data w/ missing values in black
ax = prices_interp.plot(c='r')
prices.plot(c='k', ax=ax, lw=2)
```

Visualizing the interpolated data

Using a rolling window to transform data

- Another common use of rolling windows is to transform the data

- We've already done this once, in order to *smooth* the data

- However, we can also use this to do more complex transformations

Transforming data to standardize variance

- A common transformation to apply to data is to standardize its mean and variance over time. There are many ways to do this.

- Here, we'll show how to convert your dataset so that each point represents the *% change over a previous window*.

- This makes timepoints more comparable to one another if the absolute values of data change a lot

# Transforming to percent change with Pandas

```python
def percent_change(values):
    """Calculates the % change between the last value
    and the mean of previous values"""
    # Separate the last value and all previous values into variables
    previous_values = values[:-1]
    last_value = values[-1]

    # Calculate the % difference between the last value
    # and the mean of earlier values
    percent_change = (last_value - np.mean(previous_values)) \
    / np.mean(previous_values)
    return percent_change
```

Applying this to our data

```python
# Plot the raw data
fig, axs = plt.subplots(1, 2, figsize=(10, 5))
ax = prices.plot(ax=axs[0])

# Calculate % change and plot
ax = prices.rolling(window=20).aggregate(percent_change).plot(ax=axs[1])
ax.legend_.set_visible(False)
```

Finding outliers in your data

- Outliers are datapoints that are significantly statistically different from the dataset.

- They can have negative effects on the predictive power of your model, biasing it away from its "true" value

- One solution is to *remove* or *replace* outliers with a more representative value

  Be very careful about doing this - often it is difficult to determine what is a legitimately extreme value vs an abberation

Plotting a threshold on our data

```python
fig, axs = plt.subplots(1, 2, figsize=(10, 5))
for data, ax in zip([prices, prices_perc_change], axs):
    # Calculate the mean / standard deviation for the data
    this_mean = data.mean()
    this_std = data.std()

    # Plot the data, with a window that is 3 standard deviations
    # around the mean
    data.plot(ax=ax)
    ax.axhline(this_mean + this_std * 3, ls='--', c='r')
    ax.axhline(this_mean - this_std * 3, ls='--', c='r')
```

Visualizing outlier thresholds
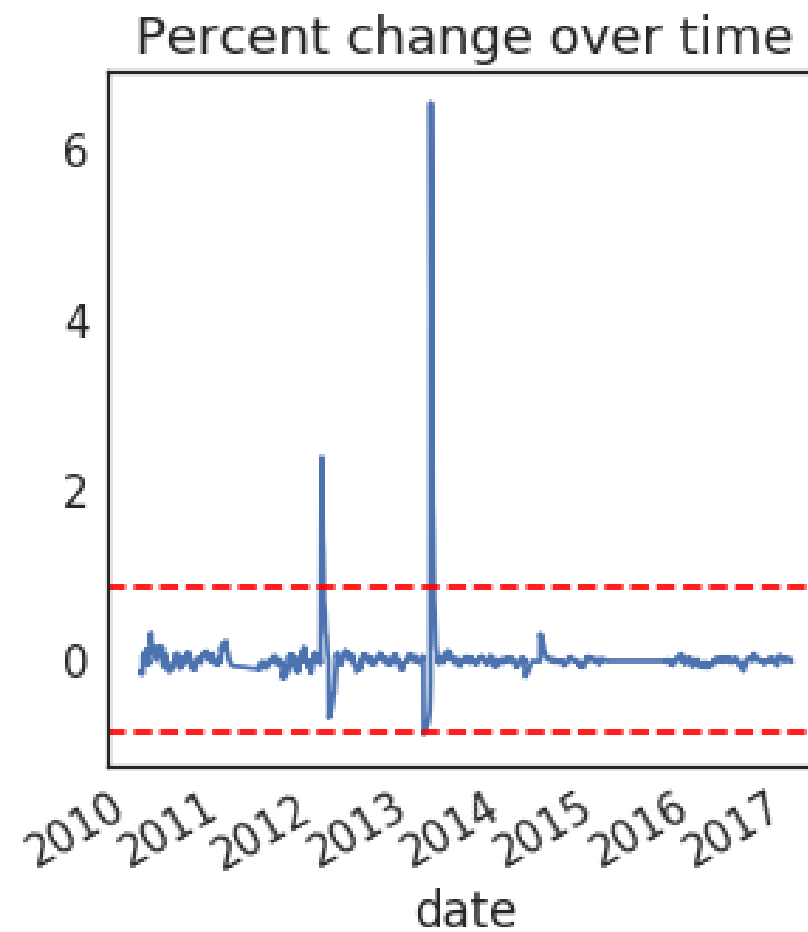
Replacing outliers using the threshold

```python
# Center the data so the mean is 0
prices_outlier_centered = prices_outlier_perc - prices_outlier_perc.mean()

# Calculate standard deviation
std = prices_outlier_perc.std()

# Use the absolute value of each datapoint
# to make it easier to find outliers
outliers = np.abs(prices_outlier_centered) > (std * 3)

# Replace outliers with the median value
# We'll use np.nanmean since there may be nans around the outliers
prices_outlier_fixed = prices_outlier_centered.copy()
prices_outlier_fixed[outliers] = np.nanmedian(prices_outlier_fixed)
```

Visualize the results

```
fig, axs = plt.subplots(1, 2, figsize=(10, 5))
prices_outlier_centered.plot(ax=axs[0])
prices_outlier_fixed.plot(ax=axs[1])
```

MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Let's practice!

MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Creating features over time

Chris Holdgraf
Fellow, Berkeley Institute for Data Science

Extracting features with windows

Raw data

func1
func2
...

Features

Using .aggregate for feature extraction

```
# Visualize the raw data
print(prices.head(3))

symbol            AIG        ABT
date
2010-01-04  29.889999  54.459951
2010-01-05  29.330000  54.019953
2010-01-06  29.139999  54.319953

# Calculate a rolling window, then extract two features
feats = prices.rolling(20).aggregate([np.std, np.max]).dropna()
print(feats.head(3))

                 AIG                   ABT
               std       amax        std       amax
date
2010-02-01  2.051966  29.889999  0.868830  56.239949
2010-02-02  2.101032  29.629999  0.869197  56.239949
2010-02-03  2.157249  29.629999  0.852509  56.239949
```

Check the properties of your features!

Using partial() in Python

```python
# If we just take the mean, it returns a single value
a = np.array([[0, 1, 2], [0, 1, 2], [0, 1, 2]])
print(np.mean(a))
1.0

# We can use the partial function to initialize np.mean
# with an axis parameter
from functools import partial
mean_over_first_axis = partial(np.mean, axis=0)

print(mean_over_first_axis(a))
[0. 1. 2.]
```

Percentiles summarize your data

- Percentiles are a useful way to get more fine-grained summaries of your data (as opposed to using `np.mean`)

- For a given dataset, the Nth percentile is the value where N% of the data is below that datapoin, and 100-N% of the data is above that datapoint.

```python
print(np.percentile(np.linspace(0, 200), q=20))
40.0
```

Combining np.percentile() with partial functions to calculate a range of percentiles

```python
data = np.linspace(0, 100)

# Create a list of functions using a list comprehension
percentile_funcs = [partial(np.percentile, q=ii) for ii in [20, 40, 60]]

# Calculate the output of each function in the same way
percentiles = [i_func(data) for i_func in percentile_funcs]
print(percentiles)
[20.0, 40.00000000000001, 60.0]

# Calculate multiple percentiles of a rolling window
data.rolling(20).aggregate(percentiles)
```

Calculating "date-based" features

- Thus far we've focused on calculating "statistical" features - these are features that correspond statistical properties of the data, like "mean", "standard deviation", etc

- However, don't forget that timeseries data often has more "human" features associated with it, like days of the week, holidays, etc.

- These features are often useful when dealing with timeseries data that spans multiple years (such as stock value over time)

datetime features using Pandas

```python
# Ensure our index is datetime
prices.index = pd.to_datetime(prices.index)

# Extract datetime features
day_of_week_num = prices.index.weekday
print(day_of_week_num[:10])
Index([0 1 2 3 4 0 1 2 3 4], dtype='object')

day_of_week = prices.index.weekday_name
print(day_of_week[:10])
Index(['Monday' 'Tuesday' 'Wednesday' 'Thursday' 'Friday' 'Monday' 'Tuesday'
 'Wednesday' 'Thursday' 'Friday'], dtype='object')
```

MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

Let's practice!