



PARALLEL COMPUTING WITH DASK

Understanding Computer Storage & Big Data

Dhavide Aruliah

Director of Training, Anaconda

"Data > one machine"



Storage units: bytes, kilobytes, megabytes, ...

watt	W	
Kilowatt	KW	10^3 W
Megawatt	MW	10^6 W
Gigawatt	GW	10^9 W
Terawatt	TW	10^{12} W

- Conventional units: factors of 1000

- Kilo \rightarrow Mega \rightarrow Giga \rightarrow Tera \rightarrow ...

Byte	B	2^3 bits
Kilobyte	KB	2^{10} Bytes
Megabyte	MB	2^{20} Bytes
Gigabyte	GB	2^{30} Bytes
Terabyte	TB	2^{40} Bytes

- Binary computers: base 2:

- Binary digit (bit)

- Byte: 2^3 bits = 8 bits

- $10^3 = 1000 \mapsto 2^{10} = 1024 \{\{4\}\}$



Hard disks



Hard disk

- Hard storage: hard disks (permanent, big, **slow**)



Random Access Memory (RAM)



RAM

- Soft storage: RAM (temporary, small, **fast**)




Time scales of storage technologies

Storage medium	Access time
RAM	120 ns
Solid-state disk	50-150 μ s
Rotational disk	1-10 ms
Internet (SF to NY)	40 ms

Storage medium	Rescaled
RAM	1 s
Solid-state disk	7-21 min
Rotational disk	2.5 hr - 1 day
Internet (SF to NY)	3.9 days

Big data in practical terms

- RAM: **fast** (ns- μ s)
- Hard disk: **slow** (μ s-ms)
- I/O (input/output) is **punitive!**



Querying Python interpreter's memory usage

```
import psutil, os

def memory_footprint():
    ...:     '''Returns memory (in MB) being used by Python process'''
    ...:     mem = psutil.Process(os.getpid()).memory_info().rss
    ...:     return (mem / 1024 ** 2)
```


Allocating memory for an array

```
import numpy as np
before = memory_footprint()

N = (1024 ** 2) // 8 # Number of floats that fill 1 MB
x = np.random.randn(50*N) # Random array filling 50 MB
after = memory_footprint()

print('Memory before: {} MB'.format(before))
```

```
Memory before: 45.68359375 MB
```

```
print('Memory after: {} MB'.format(after))
```

```
Memory after: 95.765625 MB
```

Allocating memory for a computation

```
before = memory_footprint()  
x ** 2 # Computes, but doesn't bind result to a variable
```

```
array([ 0.16344891,  0.05993282,  0.53595334, ...,  
       0.50537523,  0.48967157,  0.06905984])
```

```
after = memory_footprint()  
print('Extra memory obtained: {} MB'.format(after - before))
```

```
Extra memory obtained: 50.34375 MB
```



Querying array memory Usage

```
x.nbytes # Memory footprint in bytes (B)
```

```
52428800
```

```
x.nbytes // (1024**2) # Memory footprint in megabytes (MB)
```

```
50
```

Querying DataFrame memory usage

```
df = pd.DataFrame(x)
```

```
df.memory_usage(index=False)
```

```
0    52428800  
dtype: int64
```

```
df.memory_usage(index=False) // (1024**2)
```

```
0      50  
dtype: int64
```



PARALLEL COMPUTING WITH DASK

Let's practice!



PARALLEL COMPUTING WITH DASK

Thinking about Data in Chunks

Dhavide Aruliah

Director of Training, Anaconda

Using `pd.read_csv()` with `chunksize`

```
filename = 'NYC_taxi_2013_01.csv'

for chunk in pd.read_csv(filename, chunksize=50000):
...:     print('type: %s shape %s' %
...:           (type(chunk), chunk.shape))
```

```
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (49999, 14)
```



Examining a chunk

```
chunk.shape
```

```
(49999, 14)
```

```
chunk.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 49999 entries, 150000 to 199998  
Data columns (total 14 columns):  
medallion          49999 non-null object  
...  
dropoff_latitude   49999 non-null float64  
dtypes: float64(5), int64(3), object(6)  
memory usage: 5.3+ MB
```


Filtering a chunk

```
is_long_trip = (chunk.trip_time_in_secs > 1200)
```

```
chunk.loc[is_long_trip].shape
```

```
(5565, 14)
```

	passenger_count	trip_time_in_secs	trip_distance
167	1	300	2.1
168	3	2100	13.51
169	1	420	1.56
170	3	120	0.67
171	4	960	3.34
172	2	1140	4.13
173	5	300	2.19
174	1	1620	10.1
175	1	120	0.55
176	1	1440	10.63
177	1	120	0.47
178	1	1320	6.82
179	1	1500	5.32
180	1	420	1.71
181	3	960	4.72
182	6	1020	4.77
183	1	600	1.73

	passenger_count	trip_time_in_secs	trip_distance
168	3	2100	13.51
174	1	1620	10.1
176	1	1440	10.63
178	1	1320	6.82
179	1	1500	5.32
185	3	1260	11.17

Chunking & filtering together

```
def filter_is_long_trip(data):  
    ...:     "Returns DataFrame filtering trips longer than 20 minutes"  
    ...:     is_long_trip = (data.trip_time_in_secs > 1200)  
    ...:     return data.loc[is_long_trip]
```

```
chunks = []  
for chunk in pd.read_csv(filename, chunksize=1000):  
    ...:     chunks.append(filter_is_long_trip(chunk))
```

```
chunks = [filter_is_long_trip(chunk)  
    ...:     for chunk in pd.read_csv(filename,  
    ...:     chunksize=1000) ]
```



Using `pd.concat()`

```
len(chunks)
```


```
200
```

```
lengths = [len(chunk) for chunk in chunks]  
lengths[-5:]    # Each has ~100 rows
```

```
[115, 147, 137, 109, 119]
```

```
long_trips_df = pd.concat(chunks)  
long_trips_df.shape
```

```
(21661, 14)
```





Plotting the filtered results

```
import matplotlib.pyplot as plt
long_trips_df.plot.scatter(x='trip_time_in_secs',
                           y='trip_distance');

plt.xlabel('Trip duration [seconds]');
plt.ylabel('Trip distance [miles]');
plt.title('NYC Taxi rides over 20 minutes (2013-01-01
to 2013-01-14)');
plt.show();
```



PARALLEL COMPUTING WITH DASK

Let's practice!



PARALLEL COMPUTING WITH DASK

Managing Data with Generators

Dhavide Aruliah

Director of Training, Anaconda

Filtering in a list comprehension

```
import pandas as pd
filename = 'NYC_taxi_2013_01.csv'
def filter_is_long_trip(data):
    "Returns DataFrame filtering trips longer than 20 mins"
    is_long_trip = (data.trip_time_in_secs > 1200)
    return data.loc[is_long_trip]

chunks = [filter_is_long_trip(chunk)
           for chunk in pd.read_csv(filename,
                                     chunksize=1000)]
```




Filtering & summing with generators

```
chunks = (filter_is_long_trip(chunk)
           for chunk in pd.read_csv(filename,
                                     chunksize=1000))

distances = (chunk['trip_distance'].sum() for chunk in chunks)

sum(distances)
```

```
230909.560000000003
```



Examining consumed generators

```
distances
```

```
<generator object <genexpr> at 0x10766f9e8>
```

```
next(distances)
```

```
StopIteration          Traceback (most recent call last)  
<ipython-input-10-9995a5373b05> in <module>()
```



Reading many files

```
template = 'yellow_tripdata_2015-{:02d}.csv'

filenames = (template.format(k) for k in range(1,13)) # Generator
for fname in filenames:
    ...:     print(fname) # Examine contents
```

```
yellow_tripdata_2015-01.csv
yellow_tripdata_2015-02.csv
yellow_tripdata_2015-03.csv
yellow_tripdata_2015-04.csv
...
yellow_tripdata_2015-09.csv
yellow_tripdata_2015-10.csv
yellow_tripdata_2015-11.csv
yellow_tripdata_2015-12.csv
```



Examining a sample DataFrame

```
df = pd.read_csv('yellow_tripdata_2015-12.csv', parse_dates=[1, 2])
df.info()  # Columns deleted from output
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 71634 entries, 0 to 71633
Data columns (total 19 columns):
VendorID                71634 non-null int64
tpep_pickup_datetime    71634 non-null datetime64[ns]
tpep_dropoff_datetime   71634 non-null datetime64[ns]
passenger_count         71634 non-null int64
...
...
dtypes: datetime64[ns](2), float64(12), int64(4), object(1)
memory usage: 10.4+ MB
```

Examining a sample DataFrame

```
def count_long_trips(df):  
    ...:     df['duration'] = (df.tpep_dropoff_datetime -  
    ...:                       df.tpep_pickup_datetime).dt.seconds  
    ...:     is_long_trip = df.duration > 1200  
    ...:     result_dict = {'n_long': [sum(is_long_trip)],  
    ...:                    'n_total': [len(df)]}  
    ...:     return pd.DataFrame(result_dict)
```

Aggregating with Generators

```
def count_long_trips(df):
    ...:     df['duration'] = (df.tpep_dropoff_datetime -
    ...:                       df.tpep_pickup_datetime).dt.seconds
    ...:     is_long_trip = df.duration > 1200
    ...:     result_dict = {'n_long': [sum(is_long_trip)],
    ...:                    'n_total': [len(df)]}
    ...:     return pd.DataFrame(result_dict)

filenames = [template.format(k) for k in range(1,13)] # Listcomp

dataframes = (pd.read_csv(fname, parse_dates=[1,2])
    ...:        for fname in filenames) # Generator

totals = (count_long_trips(df) for df in dataframes) # Generator

annual_totals = sum(totals) # Consumes generators
```

Computing the fraction of long trips

```
print(annual_totals)
```

```
   n_long  n_total  
0  172617   851390
```

```
fraction = annual_totals['n_long'] / annual_totals['n_total']  
print(fraction)
```

```
0      0.202747  
dtype: float64
```



PARALLEL COMPUTING WITH DASK

Let's practice!



PARALLEL COMPUTING WITH DASK

Delaying Computation with Dask

Dhavide Aruliah

Director of Training, Anaconda

Composing functions

```
from math import sqrt
def f(z):
    ...:     return sqrt(z + 4)

def g(y):
    ...:     return y - 3

def h(x):
    ...:     return x ** 2
```

```
print(f(g(h(x)))) # Equal
```

```
x = 4
y = h(x)
z = g(y)
w = f(z)

print(w) # Final result
```

```
4.123105625617661
```

```
4.123105625617661
```



Deferring computation with delayed

```
from dask import delayed
y = delayed(h)(x)
z = delayed(g)(y)
w = delayed(f)(z)
print(w)
```

```
Delayed('f-5f9307e5-eb43-4304-877f-1df5c583c11c')
```

```
type(w) # a dask Delayed object
```

```
dask.delayed.Delayed
```

```
w.compute() # Computation occurs now
```

```
4.123105625617661
```



Visualizing a task graph

```
w.visualize()
```



Renaming decorated functions

```
f = delayed(f)
g = delayed(g)
h = delayed(h)
w = f(g(h(4)))
```

```
type(w) # a dask Delayed object
```

```
dask.delayed.Delayed
```

```
w.compute() # Computation occurs now
```

```
4.123105625617661
```



Using decorator @-notation

```
def f(x):  
    ...:     return sqrt(x + 4)  
f = delayed(f)  
  
@delayed # Equivalent to definition in above 2 cells  
...: def f(x):  
...:     return sqrt(x + 4)
```

Deferring Computation with Loops

```
@delayed
...: def increment(x):
...:     return x + 1
```

```
@delayed
...: def double(x):
...:     return 2 * x
```

```
@delayed
...: def add(x, y):
...:     return x + y
```

```
data = [1, 2, 3, 4, 5]
output = []
for x in data:
...:     a = increment(x)
...:     b = double(x)
...:     c = add(a, b)
...:     output.append(c)

total = sum(output)
```



Deferring computation with loops 2

```
total
```

```
Delayed('add-c6803f9e890c95cec8e2e3dd3c62b384')
```

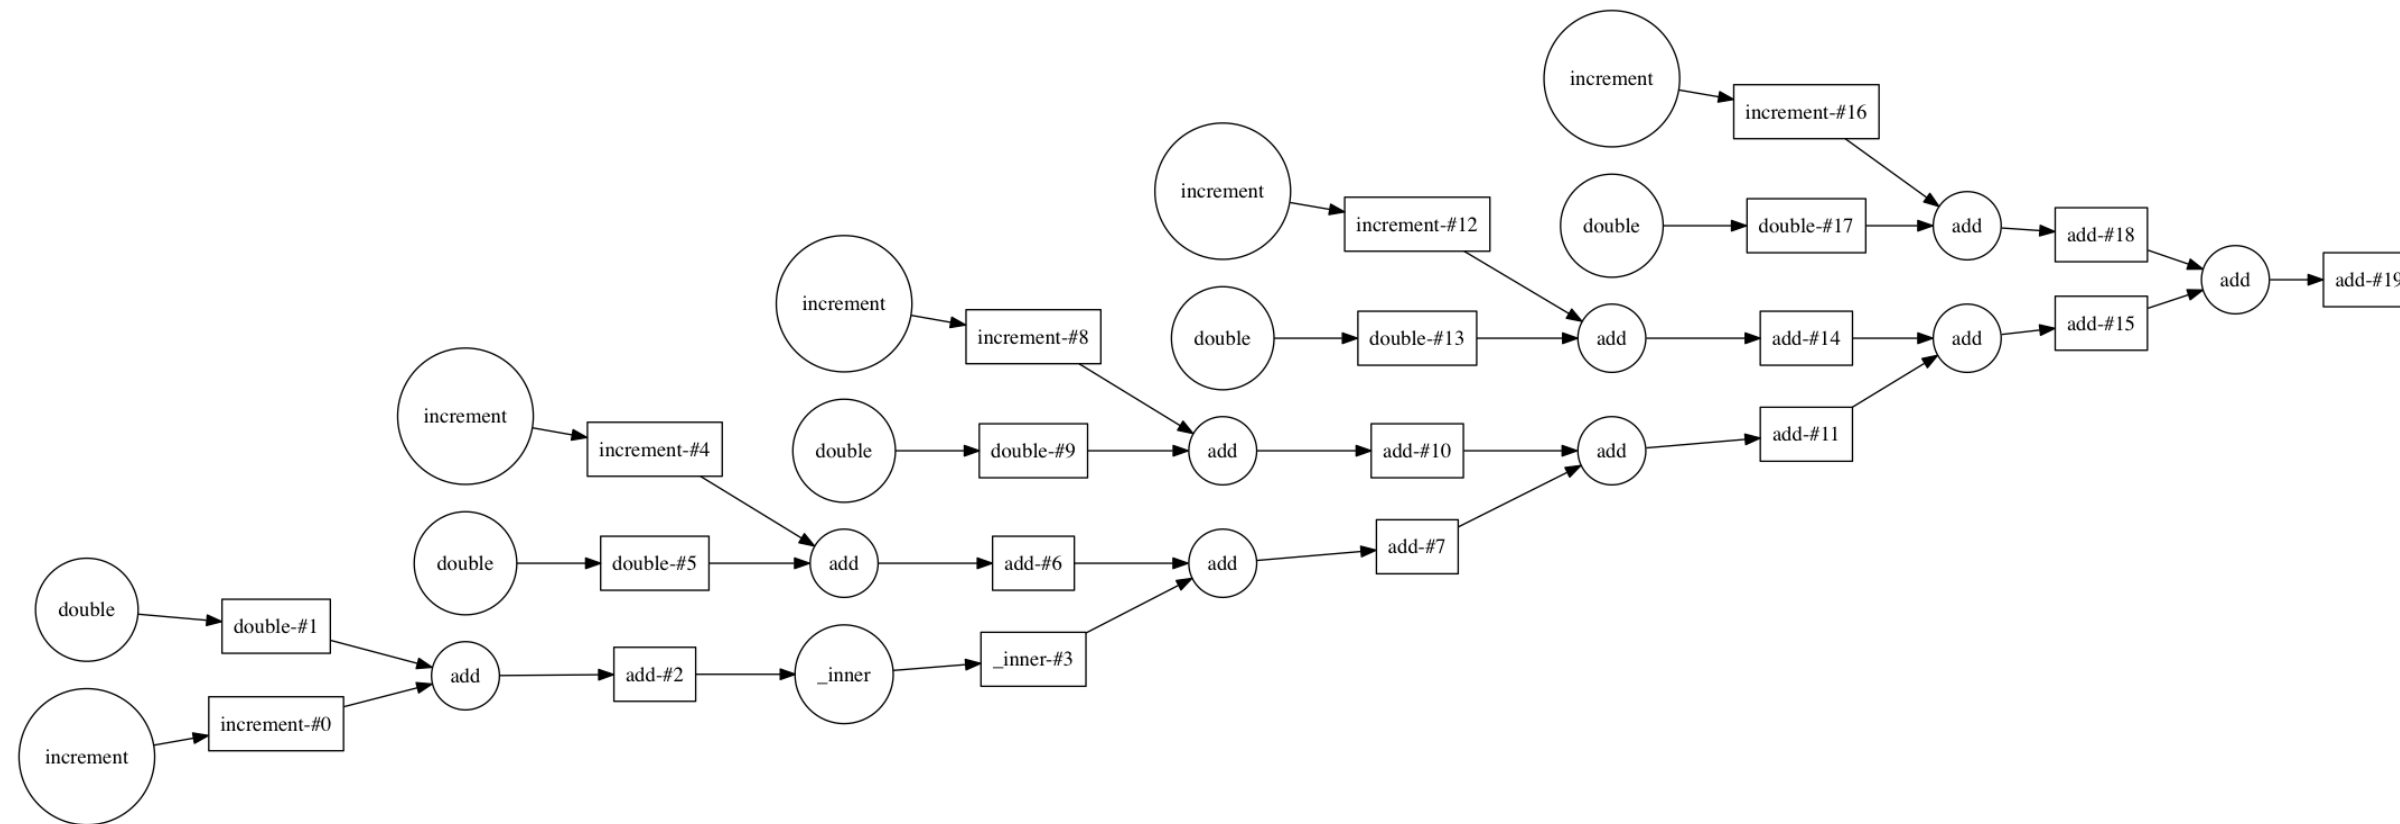
```
output
```

```
[Delayed('add-6a624d8b-8ddb-44fc-b0f0-0957064f54b7'),  
 Delayed('add-9e779958-f3a0-48c7-a558-ce47fc9899f6'),  
 Delayed('add-f3552c6f-b09d-4679-a770-a7372e2c278b'),  
 Delayed('add-ce05d7e9-42ec-4249-9fd3-61989d9a9f7d'),  
 Delayed('add-dd950ec2-c17d-4e62-a267-1dabe2101bc4')]
```

```
total.visualize()
```




Visualizing the task graph





Aggregating with delayed Functions

```
template = 'yellow_tripdata_2015-{:02d}.csv'
filenames = [template.format(k) for k in range(1,13)]

@delayed
...: def count_long_trips(df):
...:     df['duration'] = (df.tpep_dropoff_datetime -
...:                      df.tpep_pickup_datetime).dt.seconds
...:     is_long_trip = df.duration > 1200
...:     result_dict = {'n_long': [sum(is_long_trip)],
...:                   'n_total': [len(df)]}
...:     return pd.DataFrame(result_dict)

@delayed
...: def read_file(fname):
...:     return pd.read_csv(fname, parse_dates=[1,2])
```

Computing fraction of long trips with delayed functions

```
totals = [count_long_trips(read_file(fname)) for fname in filenames]

annual_totals = sum(totals)

annual_totals = annual_totals.compute()
```

```
   n_long  n_total
0  172617   851390
```

```
fraction = annual_totals['n_long']/annual_totals['n_total']
print(fraction)
```

```
0    0.202747
dtype: float64
```



PARALLEL COMPUTING WITH DASK

Let's practice!